Enterprise SLA and Capacity Metrics

Section 1: Ingestion Throughput and Worker Concurrency

The Celery worker pool processes PDF, DOCX, TXT, MD, and HTML uploads up to 50 MB per file.

Content deduplication uses SHA-256 content_hash so identical paragraphs are embedded only once per batch.

Batch database insertion commits all chunks in a single SQLAlchemy transaction after vector generation completes.

Section 2: Disaster Recovery and Database Backups

Before running any Alembic schema migration, operators execute scripts/backup_db.sh to create a timestamped pg_dump snaps

Automated pytest suites are isolated strictly to the rag_test database and abort immediately with returncode 1 if DATABASE_UR